

Patients (consented)

105 participants interviewed

Admitted for suicide risk to a psychiatric hospital in Zapopan, Mexico; aged 15-29; C-SSRS screened; interview on social-media use



80 provided an Instagram handle

81 handles: 24 not found or invalid, 14 private, 13 public without posts (78 of the 80 had completed the interview)



30 public accounts with posts

1,393 posts, 12,267 caption words; extraction 28 Aug 2026, coverage 99.8% of declared posts



18 patients with posts in both windows

Peri-episode (+/-90 d around hospital contact): 142 posts; baseline (365 d before): 157 posts



212 posts in the blind annotation kit

At most 20 per patient and window, sampled at random; 104 peri-episode, 108 baseline; opaque codes, no dates, shuffled

Comparison group (public Instagram)

3,352 profiles triaged

Seeded from school and university accounts of the same city; the patients' network (2,125 handles) and mental-health hashtags excluded a priori



123 with verifiable age evidence

Institutional tag 90, age in biography 31, age in own captions 2; no biometrics, no style-based inference



205 profiles fully extracted

122 public with posts and age evidence; 16 removed on review (5 manual, 11 by the non-person filter: organisations, brands, professionals, false age); 106 eligible



60 matched controls (2:1)

Admission envelope derived from the patients; exact strata (age band x followers x activity), nearest neighbour on metadata propensity; 18 from a neighbouring stratum

Three layers, each a pre-specified gate for the next

Layer 0: metadata

14 account-shape features (followers, rhythm, posting hour, media mix). Between persons: sampling-frame audit. Within person: posting rate.

Layer 1: writing style

25 countable features (pronouns, absolutist words, emotion lexicons, emoji valence, timing, form). Between and within persons.

Layer 2: content

8 clinical categories from a codebook; two independent language-model raters, one post per call, blind. Within person, three versions.